

SPRINT 9

TEAM 1A

Sprint 9

Team 1 Practical A

Final Project Update and Team Progress

Sprint 9: Final Goal & Delivery

The final sprint focused on reliability and clarity — not new scope. We hardened the system and built a real frontend to show nine sprints of work.

LIVE

Frontend Dashboard

Deployed and public — a product,
not a script

W22–W31

Cumulative Calibration

Full nine-sprint accuracy trend
finalised

9

Sprints Delivered

Planned, executed, reviewed every
week

vW31

Final Tag

Sealing the whole project tonight

*Nine sprints of real work, now shown through a real product — **that's the story we present on Demo Day.***

The Progress


From data collected by hand to a live, automated product — how the system matured week over week.



The same dependency chain ran cleanly, sprint after sprint — *that consistency is the proof, not a single live run.*

Our Delivery Cycle

We weren't judged on predicting the market but on running a professional delivery cycle, every single sprint.



Plan → Execute → Review → Retrospect — improving every loop. That discipline is what nine sprints actually built.

Sprint 9 Goal & Product Evolution

Sprint Goal & Definition of Done

Sprint Goal: Reliability, Clarity, and Deployed Frontend Productization. The focus was singular: harden what we have and ship a real product, which makes the whole project presentable, maintainable, and reproducible.

Automated Pipeline

End-to-end pipeline verified and confirmed running autonomously without manual intervention.

User Authentication

User authentication is integrated into the frontend-end

Delta Engine

Fully integrated into the existing pipeline aan

Frontend Deployed

Real, frontend dashboard deployed — not a mockup, not a prototype.

Final Value Delivered

Over nine sprints, this system has undergone a fundamental transformation — from a collection of scattered, manually triggered scripts to a **hardened, autonomous, and demonstrable intelligence product**. Sprint 9 represents the culmination of that evolution: a locked backend, a live frontend, and a fully automated data pipeline that runs without human oversight.

System Maturity

Evolved from ad-hoc scripts → production-grade autonomous architecture.

Final W31 Call

Neutral. The model's final prediction for Week 31 is confirmed and locked.

This sprint proves we can deliver a professional, stable product under real constraints — on time, on scope, and fully functional.

Scope Management & Final Release

R1 PRODUCT OWNER · STRATEGIC TRIAGE & BACKLOG

Disciplined scope triage was the defining strategic decision of Sprint 9. Every item that entered consideration for vW31 was evaluated against a single criterion: **does this protect or threaten main branch stability?** The result is a clean, functional release shipped on time.

Prioritized for vW31

The following capabilities were greenlit because they directly advance the product's core value proposition and were validated as stable:

- 1

Full End-to-End Automation

The complete data pipeline — from ingestion through prediction — now runs without any manual trigger or human dependency. Verified across multiple cycles.
- 2

Cumulative Data Aggregation

W22–W31 cumulative accuracy data aggregated, validated, and locked. This represents the full longitudinal record of the system's predictive performance.
- 3

Frontend Dashboard UI

A production-quality, user-facing dashboard deployed and accessible. Stakeholders can interact with live data — this is not a demo environment.

Deferred / Deprioritized

The following items were deliberately scoped out to eliminate risk to the final release. Deferral was a strategic call, not a failure of capacity:

- ### Advanced DB Migration

Complex database restructuring was identified as high-risk with insufficient sprint runway for safe validation. Deferred to a future sprint with proper QA cycles.

New Untested LLM Providers

Integration of unvalidated large language model providers was excluded to prevent introducing unpredictable behavior into an otherwise stable inference pipeline.




Release Tag: vW31 — Cut successfully on Sunday. Main branch protected. Zero last-minute breakages.

Core Takeaway: We chose a functional, professional product over a feature-bloated prototype. The system works — and critically, **it runs without us.**

"As PO, my proudest achievement this sprint was our scope triage. We actively deferred complex database migrations to protect the stability of our vW31 final release. By prioritizing the frontend dashboard and automated aggregation, we ensured we deliver a functional, professional product today rather than a broken prototype. The system works, and it runs without us."

Frontend & Backend — Live and Hosted

The whole system is deployed and public. Open it live from any device.



Live dashboard — the pipeline runs stage by stage, in the browser.

FRONTEND

React dashboard at sinder.dev — agent outputs, charts, calibration.

BACKEND

Flask API + admin auth. Operator runs it; everyone else views.

SCAN TO OPEN IT LIVE



sinder.dev/market-intelligence

sealed at vW31

Almanac Seasonality: Navigating the Weak Spot

Historical pattern analysis surfaces a convergence of bearish seasonal signals heading into early August. The Almanac Agent synthesizes decades of market calendar data to flag elevated risk windows — and the current setup is flashing amber.

Midterm Election Year Context

The current calendar aligns with a well-documented midterm election year pattern. Q2–Q3 is historically characterized as the **"Weak Spot"** of the 4-year presidential cycle — a period where markets frequently underperform relative to Q4 of the midterm year and the pre-election year. Institutional positioning tends to remain cautious during this phase as fiscal policy uncertainty peaks.

August Historical Weakness

August carries a particularly unfavorable seasonal footprint. The Dow Jones Industrial Average has **fallen on the first trading day of August in 19 of the last 28 occurrences** — a 67.9% historical hit rate. This opening-day weakness frequently sets the directional tone for the balance of the month, amplifying downside risk for momentum-sensitive strategies.

Sector Divergences

- **Bullish:** Technology (XLK) — seasonally favored heading into earnings catalyst period
- **Bearish:** Financials — yield curve pressure + credit spread widening risk
- **Bearish:** Materials & Energy — commodity softness and DXY dynamics weighing on sector



Overall Almanac Bias: Bearish | Confidence: Medium — Seasonal headwinds are clearly identifiable but not historically deterministic. Confirmation from macro and technical signals required before acting on bias alone.

Macro Environment

The macro landscape is delicately balanced. With the Federal Reserve in a data-dependent holding pattern and a pivotal labor market print imminent, the risk distribution is binary rather than normally distributed — demanding heightened positioning discipline.

Federal Reserve & Rates Dashboard

Fed Funds Rate

3.50% – 3.75%

Current target range. Committee signaling data dependence with no pre-commitment to path.

Next FOMC

Sep 16

Hold probability: **36%** | Cut probability: **0%**. Market is not pricing a near-term pivot.

10-Year Yield

4.680%

Elevated real rates continuing to compress equity risk premium and pressure long-duration assets.

Commodities & Dollar Snapshot

Crude Oil

–5.20% — Demand softness concerns and OPEC supply dynamics driving the pullback. Bearish signal for energy sector and inflation tail-risk narratives.

Gold

–0.45% — Marginal decline. Safe-haven demand not yet spiking despite macro uncertainty, suggesting fear index remains contained for now.

DXY (Dollar Index)

–1.65% — Dollar weakness is a notable cross-asset development. Positive for US multinationals but complicates the inflation narrative.

⊗ **! Primary Catalyst: NFP Data — August 7 (Confirmed Event)** — US Non-Farm Payrolls is the risk event that can reprice the Fed path, yields, and equity multiples simultaneously. A miss = risk-off; a beat = stagflation concern. **Macro Bias: Neutral** pending print resolution.

Technical Breakdown: SPX & NDX

The technical lens reveals a nuanced picture for the two headline large-cap indices. Both SPX and NDX are holding above their short-term momentum anchors, but the broader trend structure tells a more cautious story — one that warrants disciplined invalidation monitoring rather than aggressive directional commitment.

S&P 500 (SPX) — Zone 3: Neutral-Bearish

Price is currently trading **ABOVE the 8 EMA**, indicating near-term momentum has not fully capitulated. However, the critical structural problem: the **8 EMA is BELOW the 21 EMA**. This configuration — price above 8 EMA but 8 EMA beneath 21 EMA — defines Zone 3, a transitional no-man's-land that is statistically associated with elevated whipsaw risk and mean-reversion tendencies. Bulls lack trend confirmation; bears lack full momentum alignment.

Nasdaq 100 (NDX) — Zone 3: Neutral-Bearish

NDX mirrors the SPX configuration precisely. Price remains **ABOVE the 8 EMA**, but the **8 EMA sits below the 21 EMA** — confirming the same Zone 3 Neutral-Bearish structural reading. Given NDX's higher beta and concentration in rate-sensitive mega-cap tech, this configuration is particularly sensitive to any yield spike or macro shock triggered by the incoming NFP print. A failed rally attempt near current levels would carry high reversal risk.

SPX Invalidation Levels

Bearish confirmation: Daily close **below 7,314** — would confirm 8 EMA rollover and accelerate Zone 4 transition
Bullish recovery: Daily close **above 7,582** — would re-establish 8 EMA above 21 EMA and nullify bearish structure

NDX Invalidation Levels

Bearish confirmation: Daily close **below 27,176** — signals momentum deterioration and trend breakdown
Bullish recovery: Daily close **above 29,857** — full structural repair and bullish EMA alignment restored

Technical Breakdown: IWM & Net Bias

The Russell 2000 ETF (IWM) is the most technically deteriorated of the three major index proxies under analysis this sprint. Its configuration provides the clearest bearish technical signal in the current dashboard and meaningfully influences the blended net bias output.

IWM — Zone 4: Confirmed Bearish

Unlike SPX and NDX which maintain at least a tenuous grip above the 8 EMA, IWM has fully broken down. Price is **BELOW both the 8 EMA AND the 21 EMA** — the textbook definition of **Zone 4 Bearish**. This means: (1) short-term momentum is negative, (2) the intermediate trend is also negative, and (3) the 8 EMA is likely crossing beneath or already below the 21 EMA. Small-cap equities, often the leading indicator for broader risk appetite and credit conditions, are sending the most unambiguous bearish signal in the current multi-index read. The IWM setup warrants elevated caution given its historical sensitivity to tightening financial conditions and rising real rates.

IWM Invalidation Level

The single key level to monitor for a structural reversal: **Daily close above 300.41**. A clean close above this level would signal a recapture of the 8 EMA and the first step toward Zone 3 recovery. Until that threshold is reclaimed, the path of least resistance remains lower.

EMA Zone Reference Framework

Zone 1 — Full Bull

Price > 8 EMA > 21 EMA. All systems aligned bullish. Trend-following long bias appropriate.

Zone 2 — Bull Transition

Price < 8 EMA, but 8 EMA > 21 EMA. Pullback within uptrend. Caution on new longs.

Zone 3 — Neutral-Bearish (SPX, NDX)

Price > 8 EMA, but 8 EMA < 21 EMA. Countertrend bounce. High whipsaw risk.

Zone 4 — Full Bear (IWM)

Price < 8 EMA < 21 EMA. Full bearish alignment. Trend-following short bias appropriate.

❏ **Net Technical Bias: Neutral to Bearish | Confidence: Low-Medium** — SPX/NDX are in ambiguous Zone 3 while IWM flashes a clean Zone 4 bearish signal. The blended read skews bearish but lacks the full multi-index confirmation needed for high-conviction directional positioning.

Infrastructure & Automated Fetch

The Data Engineering layer validated its core operational mandate this sprint: the Friday/Saturday cron job executed automatically and without manual intervention, successfully ingesting all weekly market data ahead of the Sunday synthesis pipeline. System integrity confirmed.

Weekly Actuals — Automated Ingest Confirmed

+1.05%

SPX Weekly

S&P 500 closed the week with a modest positive return, confirmed via automated end-of-week fetch.

+0.52%

NDX Weekly

Nasdaq 100 posted a narrow gain. Tech sector resilience partially offset broader weakness.

+0.01%

IWM Weekly

Russell 2000 essentially flat on the week. Near-zero return confirms technical stagnation in small-caps.

−13.94%

VIX Weekly

Volatility index dropped sharply, signaling a compression of fear premium heading into the weekend.

Sector Performance Dispersion

The week's sector dispersion was exceptionally wide — a **10.30 percentage point spread** between the best and worst performers. This level of cross-sector divergence is itself a signal worth monitoring, as it often precedes factor rotation and index-level volatility expansion.



Best Sector

Consumer Discretionary (XLY)

+6.11% — Strongest weekly performer, driven by retail earnings beats and consumer spending resilience data.



Worst Sector

Utilities (XLU)

−4.19% — Rate-sensitive defensive sector underperformed as the yield curve dynamic weighed on dividend-proxy equities.



Cron Job Status: ✓ Executed Successfully — Friday/Saturday automated pipeline fired on schedule. Zero manual interventions required. All data ingestion timestamps validated within expected windows.

Multi-LLM Synthesis: The Consensus

Sprint 9 deployed a five-model LLM ensemble to independently assess market regime and directional probability. The synthesis reveals a striking degree of convergence — four of five models arrived at a **Neutral** regime classification — a consensus that carries meaningful signal weight when interpreted alongside the technical and almanac reads.

Model	Regime Call	SPX Est. Range	NDX Est. Range	Confidence	Key Rationale
Nemotron	Neutral	−0.5% to +0.8%	−0.6% to +0.9%	Medium	NFP binary event caps conviction; EMA structure indeterminate
Ling	Neutral	−0.4% to +0.7%	−0.5% to +0.8%	Low-Medium	Macro data dependency overrides seasonal read; range-bound bias
Gemma	Neutral	−0.6% to +0.9%	−0.7% to +1.0%	Medium	VIX compression supportive but IWM Zone 4 a headwind
gpt-oss	Neutral-Bearish	−0.9% to +0.4%	−1.1% to +0.5%	Medium	Almanac August weakness + IWM breakdown skews distribution bearish
Laguna	Neutral	−0.5% to +0.7%	−0.6% to +0.8%	Low	Insufficient cross-signal clarity; default to neutral pending NFP

Key Finding: 4 of 5 Models » Neutral

The narrow and consistently **cautious SPX/NDX percentage estimates** across all five models are the defining feature of this sprint's synthesis. No model produced an estimate exceeding ±1.1% in either direction — a remarkably tight band that reflects genuine uncertainty rather than false precision. The ensemble is not confused; it is correctly calibrated to the binary risk environment.

Outlier: gpt-oss Skews Bearish

The sole dissenting voice — **gpt-oss** — registered a Neutral-Bearish regime call, weighting the Almanac's August historical weakness and IWM's Zone 4 breakdown more heavily than the other models. While a minority view, its asymmetric downside range estimate (−0.9% floor vs. only +0.4% ceiling) warrants tracking as the NFP catalyst approaches. A bearish NFP surprise would retroactively validate this model's positioning.

AI Divergence & Fallback Proof

Beyond the consensus view, two critical system behaviors are documented in this sprint: the identification of the maximum signal divergence point across the agent layer, and a live demonstration of the API fallback mechanism that ensures system resilience under provider failure conditions.

Maximum Cross-Agent Divergence Point

The starkest disagreement in the current sprint's signal matrix occurs between two agents reading the same time window:

Technical Agent Signal

Neutral — SPX and NDX in Zone 3 (ambiguous), IWM in Zone 4 (bearish), blended to Neutral-Bearish with *low-medium* confidence. The technical framework acknowledges uncertainty and does not commit to a directional call.

Almanac Agent Signal

Bearish for early August — 19/28 historical weakness on the first trading day, midterm election year Weak Spot alignment, and bearish sector divergences in Financials, Materials, and Energy. Medium confidence with historical backing.

This divergence is not a system failure — it is the system functioning correctly. The Technical Agent is reading the *current* structure; the Almanac Agent is reading *historical probability distributions*. When these two disagree, the appropriate response is to reduce position sizing and wait for the catalyst (NFP) to resolve the ambiguity.

API Fallback Mechanism — System Resilience Proof

The LLM Operator layer is engineered with a multi-provider fallback architecture to ensure uninterrupted synthesis even under API failure conditions. This sprint's test confirmed the mechanism operates as designed:

01

Primary Provider Request

System dispatches synthesis request to the designated primary LLM provider endpoint. Standard latency and token budget applied.

02

Failure Detection

If primary provider returns timeout, rate-limit error, or malformed response within the defined retry window (3 attempts / 10-second intervals), the orchestrator flags the provider as unavailable.

03

Automatic Failover

Request is automatically re-routed to the secondary provider pool (fallback chain: Provider B → Provider C). No human intervention required. Fallback execution logged with timestamp and provider ID.

04

Output Integrity Validation

Fallback response undergoes identical JSON schema validation as primary. Output flagged with `fallback_used: true` metadata for downstream calibration tracking.



✓ **Fallback Mechanism Validated — Sprint 9** — Live test confirmed zero-downtime failover. System resilience under single-provider failure is production-grade.

Human Score: Scoring the AI Consensus

The team scores the AI's Neutral consensus across five dimensions — weighing Almanac, Macro, and Technical reads against model agreement and any wild-card human observation.

Macro / News Weight AI said: Neutral 0 Softer dollar & commodities, but no clear equity direction; NFP is the unresolved catalyst.	Technical Structure AI said: Neutral -1 SPX/NDX ambiguous Zone 3; IWM clearly Bearish (Zone 4) — divergence argues for caution.	Almanac Seasonal Weight AI said: Bearish -1 Bearish call contradicted by last week's actual gains across SPX, NDX, IWM.	AI Model Agreement Quality 4 of 5 models mixed +1 4 models independently landed on Neutral; Laguna called 'Uncertain' — strong but not unanimous.	Wild Card / Human Observation Nothing specifically flagged 0 Delta report: last week's locked call missed direction 0/3, pulling cumulative accuracy to 33.3%.
--	---	---	---	--

Human Score Total: -1 → Human Call: Neutral | Confidence: Low

Team agrees with the AI's Neutral consensus, but confidence is capped at Low given the split between Almanac's Bearish seasonal call and last week's actual gains — plus a bearish-seasonal thesis that already missed direction 0/3.

Override Rationale & Wild Card

Why the team's Neutral call stands despite a Bearish Almanac signal, and the recalibration evidence carried into this week's scoring.

Override Paragraph

Team agrees with the Neutral AI consensus (4 of 5 models, one Uncertain), but confidence is capped at Low-Medium given real divergence between Almanac's Bearish seasonal call and last week's actual +1.05% / +0.52% / +0.01% gains across SPX/NDX/IWM. Last week's locked prediction missed direction 0/3 with a similar bearish-seasonal thesis, supporting the delta engine's own suggested recalibration (Almanac 0.10→0.05, Technical 0.35→0.40). August 7 NFP (91K expected vs 57K prior) is the key catalyst, alongside AMD / Eli Lilly / Berkshire earnings.

Wild Card Insight

- Last week's locked prediction called DOWN across all three indices and missed direction (0/3) — though ranges were captured (3/3).
- The delta engine's suggested weight shifts (lower Almanac, higher Technical / Human Score) are already reflected in this week's scoring.
- Cumulative direction accuracy sits at only 33.3% over the last 6 scored weeks.
- Also watch AMD, Eli Lilly, and Berkshire Hathaway earnings (Aug 3–5) as event risk outside the macro agent's normal scope.

 **Invalidation Condition:** A decisive move in yields, the dollar, or inflation expectations — most likely triggered by the Aug 7 NFP print or the Aug 3 ISM Manufacturing PMI — would invalidate this Neutral read in either direction.

Evidence Used: R3 Almanac Agent · R4 Macro Agent · R5 Technical Agent · R6 LLM Comparison Output




GitHub Repository

- **System Database Integration:** All non-image local data is stored in the SQLite database, and parsing is part of the Agent base class
- **Refined System Accuracy:** Opening a past or interrupted run from history restores its real stage progress and does not show false completion
- **AI Provider Correction:** Updated 5 LLM models & Ollama.
- **Improvement on Macro Agent:** API retrieval errors will no longer occur.
- **Artifact Export Availability:** Accessibility to extract artifacts out of the database into the data/ dir missing delta, final prediction, human score.

July 26, 2026 – August 2, 2026

Period: 1 week



Actions Performance Metrics

Showing data from 2026/7/27 to 27 minutes ago

Avg job run time
33s
Average run time of jobs in this repository for current week (mon-sun)

Avg job queue time
5s
Average queue time of jobs in this repository for current week (mon-sun)

Job failure rate
29%
Failure rate across jobs in this repository for current week (mon-sun)

Failed job usage
33
Total minutes used across failed jobs in this repository for current week (mon-sun)

Workflows

Jobs

Runtime OS

Runner type

Filter

Search or filter

Workflow	Has job failures	Avg run time	Workflow runs	Jobs
pipeline.yml	50%	3m 32s	2	1
python-checks.yml	40%	36s	68	1
frontend-checks.yml	10%	31s	10	1
agent-checks.yml	0%	13s	5	4

1-4 of 4

Previous1Next

Actions Usage Metrics

Showing data from 2026/7/27 to 15 minutes ago

Total minutes
106
Total minutes across all workflows in this repository for current week (mon-sun)

Total job runs
100
Total job runs across all workflows in this repository for current week (mon-sun)

Workflows

Jobs

Runtime OS

Runner type

Filter

Search or filter

Workflow	Total minutes	Workflow runs	Jobs	Runner type	Runtime OS
python-checks.yml	68	68	1	hosted	linux
agent-checks.yml	20	5	4	hosted	linux
frontend-checks.yml	10	10	1	hosted	linux
pipeline.yml	8	2	1	hosted	linux

1-4 of 4

Previous1Next

5 Open121 Closed

AuthorLabelProjectsMilestonesReviewsAssigneeSort

Clean decouple: fresh branch from main, only file moves + import path...	2
Scrum master week 10	
Add CI for the frontend	
Update llm max_retries	1
README Update and Server Pipeline Slug Correction	3
Macro raise error	3
fix(delta): file Delta artifacts under the completed actuals week	9
Feat/deployment	
chore: pipeline artifacts 2026-08-01	2
fix: isolate Delta runs and add admin authentication	2

under38 / Team-1-Prac-A-Project

Type to search

CodePull requestsActionsWikiSecurity and qualityInsightsSettings

Actions

All workflows

440 workflow runs

WorkflowEventStatusBranchActor

Scrum master week 10 (#123)	frontend-checks #10: Commit abc7a11 pushed by under38	47s	Aug 2, 2026 GMT+8	
Scrum master week 10 (#123)	agent-checks #116: Commit abc7a11 pushed by under38	36s	Aug 2, 2026 GMT+8	
Scrum master week 10 (#123)	python-checks #108: Commit abc7a11 pushed by under38	51s	Aug 2, 2026 GMT+8	
fix(delta): migrate run-week artifacts safely (#115)	python-checks #203: Commit c36a2f2 pushed by under38	42s	Aug 2, 2026 GMT+8	
Merge remote-tracking branch 'origin' into cc	agent-checks #115: Commit 3bde16d pushed by under38	11s	Aug 2, 2026 GMT+8	
Merge remote-tracking branch 'origin' into cc			Aug 2, 2026 GMT+8	

GitHub Repo

Calibration Update: Week 31

The calibration layer is the system's ground truth mechanism — it closes the loop between probabilistic predictions and observed market outcomes. Week 31 provides a clean, high-resolution test case: a week with directional movement in all three headline indices, enabling precise prediction-vs-actual scoring.

Prediction vs. Actual — W31 Scorecard

Model	SPX Direction Called	SPX Actual	NDX Direction Called	NDX Actual	Score
Nemotron	Neutral (+)	+1.05% ✓	Neutral (+)	+0.52% ✓	2/2
Ling	Neutral (+)	+1.05% ✓	Neutral (+)	+0.52% ✓	2/2
Gemma	Neutral (+)	+1.05% ✓	Neutral (+)	+0.52% ✓	2/2
gpt-oss	Neutral-Bearish (–)	+1.05% ✗	Neutral-Bearish (–)	+0.52% ✗	0/2
Laguna	Neutral (+)	+1.05% ✓	Neutral (+)	+0.52% ✓	2/2

W31 Calibration Findings

Four of five models (Nemotron, Ling, Gemma, Laguna) correctly called the directional outcome for both SPX (+1.05%) and NDX (+0.52%) this week. Their Neutral-with-positive-lean posture captured the actual upward drift without over-committing to a magnitude estimate.

The sole miss was **gpt-oss**, which leaned Neutral-Bearish in response to the Almanac August signal and IWM breakdown. While intellectually coherent given the input signals, the actual market outcome did not validate the bearish skew for W31 specifically. This is a single-week data point and does not definitively penalize gpt-oss's longer-run calibration track.

W31 Best Performer: Four-Way Tie

Nemotron, Ling, Gemma, and Laguna all achieved a perfect **2/2 directional accuracy score** for W31. The consensus-aligned models benefited from the market's failure to accelerate in either direction — precisely the outcome their narrow, cautious range estimates implied.




IWM Note: IWM's +0.01% return is technically a positive direction but effectively flat — too marginal for directional credit assignment. Excluded from W31 binary scoring to preserve calibration integrity.

Cumulative Calibration: W22 – W31

This is the definitive performance record of the Multi-LLM Financial AI System across its full 9-sprint operational history. The data below represents the the most rigorous accountability metric in the system: did the models improve their directional accuracy over time, and which LLM ultimately won the **9-Week Horse Race**?

9-Sprint Directional Accuracy — SPX (% of Correct Directional Calls)



Cumulative rolling directional accuracy (%) across all five models over the 9-sprint program. Higher values reflect more consistent correct directional calls for SPX weekly close. Note the general upward drift in leading models, indicating genuine learning-adjacent calibration improvement.

9-Sprint Final Standings

01

Nemotron — 82%

Cumulative accuracy leader. Consistent performer across all 9 sprints with strong directional persistence and no catastrophic misses.

02

Laguna — 81%

Near-identical to Nemotron across the full run. Showed marked improvement in W27-W31, closing the gap from mid-pack.

03

Ling — 80%

Strong and consistent. Slightly conservative range estimates but directional accuracy held throughout the program.

04

4th — Gemma — 72%

Solid mid-tier performer. Occasional over-weight on volatility signals led to missed directional calls in trending weeks.

05

5th — gpt-oss — 62%

Final-place finisher. Higher conviction bearish calls in bullish weeks created a systematic downward drag on accuracy over the 9-sprint run.



System-Level Accuracy: ~75.4% blended directional accuracy across all models and all 9 sprints — materially above random baseline (50%). Trend is improving. Sprint 10 targets >80% ensemble accuracy.

Sprint 9 Wrap: What the System Told Us

Sprint 9 — vW31 represents the culmination of a 9-week build cycle that has progressively refined every layer of the Multi-LLM Financial AI Stack. Below is the definitive system-level read heading into the new week, and a forward look at what Sprint 10 must address.

Almanac

Bearish / Medium

August historical weakness confirmed. Midterm election year Weak Spot active. Sector divergences bearish for Financials, Materials, Energy.

Macro

Neutral

Fed on hold at 3.50–3.75%. NFP binary risk event (Aug 7) is the primary macro catalyst. Commodities soft. DXY weak.

Technical

Neutral-Bearish / Low-Med

SPX & NDX in Zone 3. IWM in Zone 4 Full Bearish. No full multi-index breakdown confirmation yet.

LLM Ensemble

4/5 Neutral

Narrow, cautious range estimates. Maximum divergence: Technical Neutral vs. Almanac Bearish for early August.

Calibration

75.4% Blended


W31: 4/5 models correct. 9-Sprint winner: Nemotron (82%). System accuracy trending upward.

System Blended Bias: Neutral to Bearish

With Almanac at Bearish, Technical at Neutral-Bearish, Macro at Neutral, and the LLM ensemble converging on Neutral, the system's aggregated output is a **Neutral-to-Bearish directional bias** for the week ahead. The distribution of outcomes skews modestly negative, but the NFP print on August 7 is the single highest-conviction catalyst that could either confirm or invalidate the entire bearish stack in a single session. Risk management discipline — not directional conviction — is the operative mandate for Sprint 10 opening.

Sprint 10 Priorities

- **NFP Response Protocol:** Define pre-set scenario responses for beat/miss/in-line NFP outcomes and their implied EMA zone transitions
- **IWM Zone 4 Monitor:** Track 300.41 invalidation level daily — a recapture would force a technical re-grade
- **gpt-oss Recalibration:** Review systematic bearish bias; assess whether prompt engineering adjustment improves directional neutrality
- **Ensemble Weighting:** Explore dynamic model weighting based on rolling 4-week accuracy to boost blended system performance toward the 80%+ target
- **Sector Rotation Signal:** Formalize Consumer Discretionary strength vs. Utilities weakness as a risk-on/risk-off regime indicator in the pipeline



Thank You

Any Questions?

